

AgroVoice-XAI: Smart Visual Crop Advisory System

System Architecture, Deep Learning Model Selection, Explainable AI (Grad-CAM), Microclimatic Risk Engine, and Multilingual Agronomic Voice Pipeline

PROJECT ENGINEERING TEAM

Shashwat Shaurya

USN: 1MV21CS098 | shashwat.shaurya2005@gmail.com

Parth Verma

USN: 1MV21CS075 | vermaparth3108@gmail.com

Shwetanshu

USN: 1MV21CS107 | shwetanshu44@gmail.com

Kumar Vaibhav

USN: 1MV21CS054 | k.vaibhav10201@gmail.com

Project Guide:

Dr. Pallavi, Associate Professor

Head of Department:

Dr. Anitha T N, HOD, Dept. of CSE

Principal:

Dr. M N Thippeswamy, Principal

Department of Computer Science and Engineering

Sir M. Visvesvaraya Institute of Technology (Sir MVIT), Bengaluru – 562157

Affiliated to Visvesvaraya Technological University (VTU), Belagavi, Karnataka, India

1. Executive Overview & Problem Formulation

Agriculture represents the economic backbone of rural India, employing over 58% of the rural workforce. However, crop yield loss due to infectious phytopathological agents (fungi, bacteria, viruses, oomycetes, and acari) accounts for an estimated ₹1.5 lakh crore (\$18+ billion USD) annually. Smallholder farmers face three critical systemic bottlenecks:

- 1. **The Visual Diagnosis Fallacy (Black-Box AI):** Traditional deep learning image classifiers output a probability distribution across classes (e.g. Tomato___Early_blight: 92%) without visual grounding. Farmers and agronomists cannot verify whether the model recognized actual pathogenic lesion necrosis or took opportunistic "shortcut learning" from background soil, fingers, or lighting conditions.
- 2. **Meteorological Disconnect:** Plant disease proliferation is strictly governed by the *Disease Triangle* (Susceptible Host + Virulent Pathogen + Conducive Environment). An image-only classifier operates in isolation from environmental physics, ignoring that a fungal spore requires leaf wetness ($\geq 4 \text{ ext}\{ \text{hours}\}$) and high relative humidity ($\geq 75\%$) to germinate.
- 3. **Linguistic & Technical Inaccessibility:** Academic prescriptions filled with complex Latin binomials (e.g. *Venturia inaequalis*, *Cercospora sojina*) and chemical jargon without localized vernacular translations (Kannada, Hindi) or oral speech synthesis fail to serve farmers with basic formal education.

AgroVoice-XAI solves these challenges by uniting **Edge Deep Learning (MobileNetV2)**, **Explainable AI Visual Saliency (Grad-CAM)**, **Real-time Microclimatic Risk Modeling (Open-Meteo API)**, a **4-Pillar Clinical Agronomic Knowledge Base (39 Taxa)**, and a **Bidirectional Indic Voice Assistant (Web Speech API)** in a single, responsive, daylight-optimized web application.

2. End-to-End Multi-Tier System Architecture

The platform follows a decoupled, resilient, service-oriented multi-tier architecture designed for low-latency inference (<450ms) and field reliability.

Tier 1: Client / Edge Presentation Layer (React 18 + Vite)

Port: 5173 / HTTPS

Responsive Single Page Application (SPA) providing live camera stream capture, hardware torch toggle, front/rear camera selection, drag-and-drop file ingestion, Grad-CAM side-by-side visualizer, 4-pillar advisory tabs, and Web Speech STT/TTS audio synthesis engine.

Tier 2: Asynchronous Application & Routing Layer (FastAPI)

Port: 8000 / ASGI Uvicorn

High-throughput asynchronous Python ASGI server exposing RESTful endpoints (/api/analyze-crop, /api/weather-context, /api/voice-query, /api/health). Manages multipart image upload buffering, pipeline orchestration, and response schema serialization via Pydantic.

Tier 3: AI Inference & Visual Explainability Engine (PyTorch + Grad-CAM)

Model: MobileNetV2 (39 Classes)

Executes deep convolutional forward pass on preprocessed $224 \times 224 \times 3$ image tensors. PyTorch forward and full backward hooks intercept layer activations and gradients from features[-1] to synthesize class-discriminative heatmaps blended via OpenCV Jet colormap.

Tier 4: Environmental Risk & Microclimate Engine (Open-Meteo Integration)

WMO Grid Model (0.1° Resolution)

Asynchronously queries real-time hourly meteorological telemetry (Temperature, Relative Humidity, Wind Speed, Precipitation) based on client GPS coordinates. Computes continuous risk scores $(0-100)$ against biological pathogen threshold envelopes.

Tier 5: 4-Pillar Clinical Agronomic Knowledge Base & Indic NLP JSON Taxonomy + Regex Intent Parser

Stores verified 4-pillar treatment regimens (Etiology, Organic/Biological, Chemical with exact active ingredients & 15L pump dosages, Cultural Prevention) translated into native Kannada, Hindi, and English.

3. Deep Learning Model Selection & Comparative Rationale

3.1 Why MobileNetV2 was Chosen

For field agricultural deployment, model selection is governed by a strict trade-off between **top-1 classification accuracy**, **parameter size**, **inference latency on resource-constrained hardware**, and **computational efficiency (FLOPs)**.

Model Architecture	Parameters (M)	FLOPs (G)	Top-1 Val Accuracy	Inference Latency (CPU)	Edge Readiness / Suitability
MobileNetV2 (Our Selection)	3.5M	0.30G	98.43%	22 ms	Optimal for Farm Mobile & Edge
ResNet-50	25.6M	4.10G	97.80%	89 ms	Heavy memory footprint; slow on edge
VGG-16	138.4M	15.50G	95.20%	240 ms	Prohibitive storage & energy consumption
EfficientNet-B0	5.3M	0.39G	98.10%	38 ms	High accuracy but complex kernel ops
Vision Transformer (ViT-B/16)	86.6M	17.60G	96.90%	310 ms	Requires huge datasets; poor on CPU

3.2 Architectural Highlights of MobileNetV2

MobileNetV2 introduces two fundamental architectural innovations that allow it to match or exceed the accuracy of ResNet-50 while utilizing **7 times fewer parameters**:

1. **Depthwise Separable Convolutions:** Factorizes standard convolution into:
 - Depthwise Convolution:* Applies a single 3×3 spatial filter per input channel.
 - Pointwise Convolution:* Applies a 1×1 convolution to linearly combine the depthwise outputs across channels.

This reduces computational cost by a factor of:

$$\text{Reduction Ratio} = \frac{D_K \cdot D_K \cdot M \cdot D_F \cdot D_F + M \cdot N \cdot D_F \cdot D_F}{D_K \cdot D_K \cdot M \cdot N \cdot D_F \cdot D_F} = \frac{1}{N} + \frac{1}{D_K^2}$$

$\approx \frac{1}{9}$ for 3×3 kernels

2. **Inverted Residual with Linear Bottleneck:** Unlike classic ResNet blocks which compress channels then expand them, MobileNetV2 uses an *Inverted Residual Block* that first expands input channels by an expansion factor $t=6$, performs depthwise filtering in the high-dimensional manifold, and then projects

back to a lower-dimensional bottleneck using a *linear activation* (preventing ReLU from destroying non-linear features in low-dimensional space).

3.3 Dataset, Taxonomy & Training Specifications

- **Combined Agricultural Dataset:** 39 distinct classes (Apple, Blueberry, Cherry, Corn, Grape, Orange, Peach, Pepper, Potato, Raspberry, Soybean, Squash, Strawberry, Tomato, and Background/Non-leaf).
- **Total Dataset Volume:** 57,811 high-resolution images stratified into Train (85%) and Validation (15%).
- **Hardware Acceleration:** Trained on an NVIDIA GeForce RTX 4060 Laptop GPU (CUDA 12.x, 8GB VRAM).
- **Optimization & Hyperparameters:** Adam optimizer ($\beta_1=0.9$, $\beta_2=0.999$), weight decay 1×10^{-4} , Cross-Entropy Loss, learning rate 1×10^{-4} with StepLR decay, batch size 64.
- **Achieved Validation Metrics:** **98.43% Validation Accuracy**, **98.81% Training Accuracy**, Cross-Entropy Loss of **0.0354**.

4. Explainable AI (XAI) with Grad-CAM

4.1 Theoretical Foundation of Grad-CAM

Gradient-weighted Class Activation Mapping (Grad-CAM) generates visual explanations for decisions made by convolutional neural networks without requiring structural retraining. In our architecture, hooks are attached to the final convolutional layer of MobileNetV2 (features[-1]), producing a feature activation map $A \in \mathbb{R}^{1280 \times 7 \times 7}$.

To calculate the importance of each feature map k for predicted disease class c , we compute the gradient of the winning class score y^c (before softmax) with respect to feature map activations A^k :

$$\alpha_k^c = \frac{1}{Z} \sum_{i=1}^U \sum_{j=1}^V \frac{\partial \text{Global Average Pooling}}{\partial A_{i,j}^k}$$

The class-discriminative localization map $L_{\text{Grad-CAM}}^c$ is computed as the rectified linear combination of forward activation maps weighted by α_k^c :

$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left(\sum_k \alpha_k^c A^k \right)$$

The ReLU function ensures that only features that have a positive influence on the target disease class are retained, filtering out features that belong to other classes or background clutter.

4.2 Pipeline Implementation & Color Blending

1. **Gradient Capture:** A forward hook stores A^k , and a backward hook captures $\frac{\partial y^c}{\partial A^k}$ during backpropagation.
2. **Min-Max Normalization:** $L_{\text{Grad-CAM}}^c$ is normalized to $[0.0, 1.0]$.
3. **Bilinear Upsampling:** Saliency map is resized from 7×7 to the original image dimensions ($W \times H$).
4. **Colormap Transformation:** Saliency is converted to a 3-channel RGB Jet Colormap (where red denotes highest neural activation and blue denotes low activation).
5. **Alpha Blending:** The heatmap is blended with the original leaf photo using $\alpha = 0.45$:

$$I_{\text{Blended}} = (1 - \alpha) \cdot I_{\text{Original}} + \alpha \cdot I_{\text{Heatmap}}$$

6. **Base64 Data URI Serialization:** Output image is encoded into a high-efficiency PNG Data URI string returned in the JSON payload to the client.

5. Environmental Risk & Microclimatic Triangulation Engine

Plant pathogens require specific environmental conditions to complete their biological life cycle. The **Environmental Risk Engine** connects live meteorological sensor telemetry with biological pathogen thresholds:

$$\text{Risk Score } S_{\text{Risk}} = S_{\text{Baseline}} + \Delta S_{\text{Humidity}} + \Delta S_{\text{Temperature}} + \Delta S_{\text{Rain}} + \Delta S_{\text{Wind}} \quad (\text{Bounded } [0, 100])$$

- **S_{Baseline} (Pathogen Biological Virulence):** Critical severity pathogens (e.g. Late Blight) start with an inherent baseline of \$+40\$, Moderate with \$+20\$, Healthy with \$0\$.
- **$\Delta S_{\text{Humidity}}$ (Spore Germination):** If live relative humidity $RH \geq RH_{\text{min}}$ (typically 75%), adds \$+25\$ to \$+30\$ points due to surface moisture film formation.
- **$\Delta S_{\text{Temperature}}$ (Thermal Incubation):** If live temperature falls within the optimal pathogen replication range $[T_{\text{min}}, T_{\text{max}}]$ (e.g. 18°C - 28°C), adds \$+20\$ points.
- **ΔS_{Rain} (Foliar Wetness):** Recent rainfall (>0 mm) prolongs leaf wetness, adding \$+15\$ points.
- **ΔS_{Wind} (Airborne Spore Dispersion):** Wind speeds >15 km/h facilitate airborne spore dispersal across crop rows, adding \$+10\$ points.

Risk Score Range	Risk Level	Agronomic Meaning & Farmer Action
\$0 - 25\$	Low Risk (●)	Weather does not favor spread; maintain routine field scouting and sanitary hygiene.
\$26 - 50\$	Moderate Risk (●)	Microclimate supports localized infection; apply organic neem or biocontrol spray.
\$51 - 75\$	High Risk (●)	Pathogen actively spreading; prepare preventative fungicide / bactericide spray.
\$76 - 100\$	Critical Risk (●)	Severe epidemic risk; immediately spray curative fungicides and isolate infected areas.

6. 4-Pillar Clinical Agronomic Knowledge Engine

Unlike generic text outputs, the system structures every diagnosis into four standardized, actionable pillars designed in consultation with agricultural extension protocols:

Pillar 1: Etiology & Root Cause (ರೋಗಕ್ಕೆ ಕಾರಣ / बीमारी का कारण)

Explains the exact biological and environmental trigger in plain language (e.g. "Whiteflies sucked sap from leaf veins, transmitting the curl virus under hot, dry conditions").

Pillar 2: Organic & Biological Remedy (ಮನೆಮದ್ದು & ಸಾವಯವ ಪರಿಹಾರ / जैविक उपचार)

Provides direct, eco-friendly recipes with standard farm measurements (e.g. "Mix 5 ml Neem Oil + 2 ml soap in 1 liter of water" or "Spray 100 ml fresh Gomutra + 5g Trichoderma in 1L water").

Pillar 3: Chemical Spray & Exact Pump Dosage (ಔಷಧಿ ಸಿಂಪರಣೆ / दवा छिड़काव)

Specifies common commercial brand names and active ingredients with exact dosages for standard **15-liter backpack spray pumps** (ಬೆನ್ನು ಪಂಪ್ / 15L टंकी) and Pre-Harvest Intervals (PHI).

Pillar 4: Cultural Preventive Measures (ತಡೆಗಟ್ಟುವ ಕ್ರಮಗಳು / रोकथाम)

Clear field sanitation steps to protect future crop cycles (crop rotation, drainage channels, eradication of infected debris, weed borders).

7. Multilingual Indic Voice & Accessibility Pipeline

To ensure full usability for farmers of all literacy levels, the platform embeds a browser-native, zero-latency speech pipeline supporting ****Kannada (`kn-IN`)****, ****Hindi (`hi-IN`)****, and ****English (`en-IN`)****:

- **Speech-to-Text (STT):** Utilizes the `webkitSpeechRecognition` engine configured for continuous acoustic stream capture and localized Indic phoneme decoding.
- **Intent Classification:** Regex-based semantic parser maps farmer vocal questions into specific intents (`ask_treatment`, `ask_organic`, `ask_chemical`, `ask_prevention`, `ask_cause`, `ask_weather`).
- **Text-to-Speech (TTS):** Utilizes the `SpeechSynthesisUtterance` API with calibrated pitch (\$1.0\$), rate (\$0.92\$), and regional accent models for clear, natural audio advisory playback.

8. Complete Technology Stack Matrix

Layer / Component	Technology / Framework	Exact Version	Functional Role in System
Frontend UI Framework	React + Vite	18.3.1 / 5.4.21	Reactive component rendering, state management & instant hot reloading
Styling & Design System	TailwindCSS + Lucide Icons	3.4.17 / 0.460.0	Daylight-optimized agricultural palette, card layouts & intuitive iconography
Speech Processing	Web Speech API (STT/TTS)	W3C Standard	Native multilingual speech recognition and voice synthesis in browser
Backend API Framework	FastAPI + Uvicorn	0.115.6 / 0.34.0	Asynchronous REST API routing, CORS handling & multipart image parsing
Deep Learning Engine	PyTorch + Torchvision	2.6.0+cu124	Tensor computations, MobileNetV2 backbone & gradient backpropagation
Computer Vision & XAI	OpenCV + Pillow + NumPy	4.11.0 / 11.1.0	Spatial resizing, Jet colormap transformation, alpha blending & DataURI encoding
Meteorological Data	Open-Meteo REST API	v1 (Free/Open)	Hourly microclimatic temperature, humidity, wind & precipitation telemetry
Hardware Acceleration	NVIDIA GeForce RTX 4060	CUDA 12.4	GPU model training, weight tensor optimization and batch evaluation

9. Conclusion & Future Roadmap

AgroVoice-XAI successfully bridges the gap between state-of-the-art computer vision and real-world farm decision making. By augmenting high-accuracy MobileNetV2 classification (98.43%) with visual Grad-CAM grounding, live microclimatic risk triangulation, actionable 4-pillar clinical prescriptions, and multilingual voice capabilities in Kannada and Hindi, the platform provides smallholder farmers with an indispensable, trustworthy, and accessible AI plant doctor.

Future Enhancements:

- Edge On-Device Inference via ONNX Runtime Web:** Compiling MobileNetV2 to WebAssembly/WebGPU for zero-internet 100% offline edge diagnosis on smartphones in remote fields.
- Drone Multispectral Imagery Ingestion:** Extending the FastAPI ingestion layer to accept orthomosaic multispectral drone maps for automated farm-wide disease heatmapping.
- Localized Market & Pesticide Price Aggregation:** Connecting real-time KVK / APMC market feeds to show current local prices of recommended organic and chemical remedies.